



Technische Hochschule
Ingolstadt

Principles of Autonomy and Decision Making

(AI_PrincAutonomy_2808)

Week 5: Value- and Policy-Iteration

Team:

- Prof. Dr. rer. nat. Lenz Belzner
- Chidvilas Karpenahalli Ramakrishna, M.Eng.
- Adithya Mohan, M.Eng.
- Zahra Zeinaly, Ph.D.

MDPs: Recap



- MDP components:
 - States $s \in S$
 - Actions $a \in A$
 - Reward function $R(s, a, s')$
 - Discount factor γ
 - Transition function $T(s, a, s') = P(s'|s, a)$
- Quantities:
 - Optimal Policy
 - $\pi^*(s)$
 - Maps states to actions
 - Utility
 - Expected sum of discounted rewards
 - Optimal value of a state
 - $V^*(s)$
 - Expected future utility from a state and acting optimally thereafter
 - Optimal Q-value of a state-action pair
 - $Q^*(s, a)$
 - Expected future utility of a state by taking an action a and acting optimally thereafter

Bellman Equation: Recap



- Bellman equation:
 - Links the quantities and shows what is needed to take optimal actions
 - Take correct first action
 - Keep being optimal later
- Derivation:
 - $V^*(s) = \max_a Q^*(s, a)$
 - $Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
 - Bellman equation:
 - $V^*(s) = \max_a \sum_{s'} P(s'|a, s) [R(s, a, s') + \gamma V^*(s')]$
- How to solve the Bellman equation?
 - Value-iteration
 - Policy-iteration
 - Dynamic programming (both value- and policy-iteration are DP)
 - Reinforcement Learning




Fig 1. MDP search tree

Week 5: Value- and Policy-Iteration

Value-Iteration: Introduction



- Aim:
 - We wish to find optimal values $V^*(s)$ of each state using which an optimal policy $\pi^*(s)$ can be extracted.
- $V^*(s) = \max_a \sum_{s'} P(s'|a, s) [R(s, a, s') + \gamma V^*(s')]$
 - Here, $V^*(s)$ is the value of being in a state and acting optimally
 - How to obtain $V^*(s)$?
 - Value-iteration
- Value-iteration:
 - Update values for k iterations
 - We keep $V_0(s)=0$
 - We define $V_k(s)$ is the value of a state in k^{th} iteration
 - $V_k(s) = \max_a \sum_{s'} P(s'|a, s) [R(s, a, s') + \gamma V_{k-1}(s')]$
 - Stop when $\max_s |V_k(s) - V_{k-1}(s)| < \epsilon$ i.e. convergence




Fig 1. Value-iteration search tree

Value-Iteration: Policy Extraction



- Extracting optimal policy $\pi^*(s)$ from $V^*(s)$ of each state
 - $\pi^*(s) = \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$
 - $\pi^*(s) = \underset{a}{\operatorname{argmax}} Q^*(s, a)$
- Demo:
 - Week 5 lecture code – *value_iteration.py*
 - *NOTE: See how policy converged much earlier than values*

Value-Iteration: Problems



- Value-iteration is slow:
 - Why?
 - $V_k(s) = \max_a \sum_{s'} P(s'|a, s) [R(s, a, s') + \gamma V_{k-1}(s')]$
 - \max_a is very slow
- The values of each state do not change significantly after certain iterations
- The policy often converges before the values
 - Wasted computation
- Solution:
 - Occasionally check if the policy has converged and stop computation if convergence is reached
 - Policy-iteration



-

Fig 1. Policy-iteration search tree

Policy-Iteration



- Steps:
 1. Policy evaluation:
 - Fix a policy $\pi(s)$
 - Get utilities i.e. $V^\pi(s)$ for $\pi(s)$ until convergence (not $V^*(s)$)
 2. Policy improvement:
 - Update policy $\pi(s)$ using one-step look-ahead with resulting converged utilities $V^\pi(s)$
 3. Iteration:
 - Repeat steps 1 and 2 until policy converges
- Advantage:
 - We check $\pi(s)$ for convergence before continuing the iteration
 - This might result in faster convergence of $\pi(s)$
- NOTE:
 - $\pi(s)$ - Deterministic policy
 - $\pi(a|s)$ - Stochastic policy




Fig 1. Policy-iteration search tree

Policy-Iteration: Policy Evaluation



- Consider a fixed i^{th} policy $\pi^i(s)$
- Get utilities:
 - Update the value of each state using the fixed policy to choose actions
 - $V_k^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i, s') + \gamma V_{k-1}^{\pi_i}(s')]$
 - Continue for a fixed number of iterations m
 - Upon completing m iterations, update the policy to get $\pi^{i+1}(s)$, check for convergence, if not, repeat
 - NOTE: We do not have \max_a in this equation which reduces computation




Fig 1. Policy-iteration search tree

Policy-Iteration: Policy Improvement



- Update policy:
 - $\pi_{i+1}(s) = \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$
 - We updated our old policy $\pi^i(s)$ to $\pi^{i+1}(s)$
 - Upon updating the policy repeat step 1 (Policy evaluation) for another m iterations
 - Follow this until policy converges i.e. $\pi_{i+n}(s) = \pi_{i+n-1}(s)$
- Demo:
 - Week 5 lecture code – `policy_iteration.py`
 - *NOTE: See how policy-iteration ran much faster than value-iteration*




Fig 1. Policy-iteration search tree

Summary



- Value-iteration:
 - We update the utilities till convergence to get the optimal values $V^*(s)$
 - We use $V^*(s)$ to extract an optimal policy $\pi^*(s)$ at the end
 - Slower convergence as policy might converge before values
- Policy-iteration:
 - We update utilities for m number of iterations using a fixed policy $\pi_i(s)$
 - After reaching m iterations, I update my policy to $\pi_{i+1}(s)$
 - I repeat the above steps until policy converges
 - Faster convergence than value-iteration
- Motivation for Dynamic programming (DP):
 - DP is a powerful computational technique that solves problems by breaking them down into simpler subproblems, *solving each of these subproblems just once, and storing their solutions*
 - Both value-iteration and policy-iteration are dynamic programming techniques