



Technische Hochschule
Ingolstadt

Principles of Autonomy and Decision Making

(AI_PrincAutonomy_2808)

Week 8: Reinforcement Learning

Guest Lecturer: Zahra Zeinaly, Ph.D.

Ref: Adapted from RL Course by David Silver

Team:


- Prof. Dr. rer. nat. Lenz Belzner
- Chidvilas Karpenahalli Ramakrishna, M.Eng.
- Adithya Mohan, M.Eng.
- Zahra Zeinaly, Ph.D.

Contents



- About Reinforcement Learning problem
- The Reinforcement Learning Formalism
- Inside an RL agent
- Problems within Reinforcement Learning

Many Faces of Reinforcement learning




Ref: RL course by David Silver

Reinforcement Learning in a nutshell

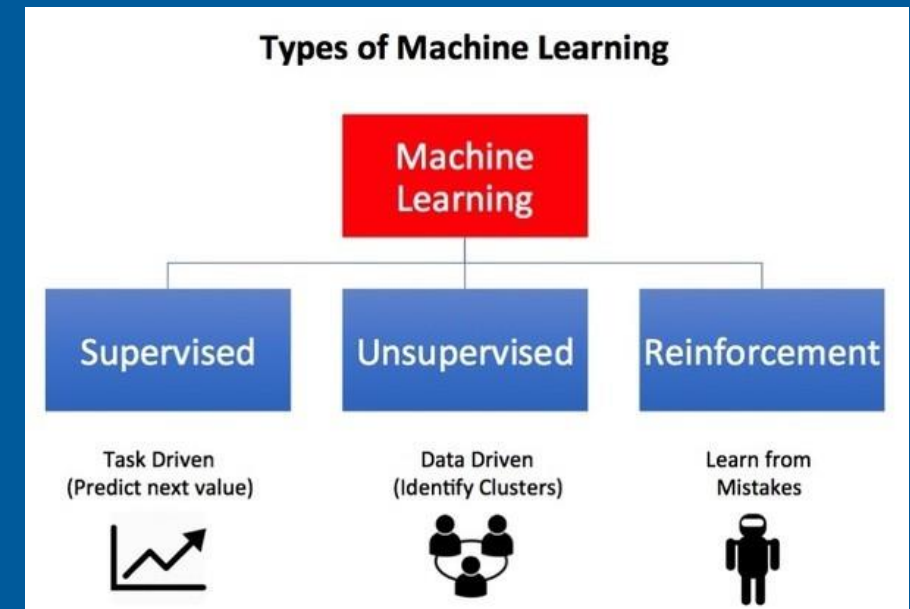


- RL is a general-purpose framework for decision-making
 - RL is for an **agent** with the capacity to **act**
 - Each **action** influences the agent's future **state**
 - Success is measured by a scalar **reward** signal
 - Goal: **select actions to maximise future reward**
 - Learning rather than direct planning



What makes reinforcement learning different from other machine learning paradigms?

- There is no supervisor, only a reward signal.
- The agent learns by interacting with environment.
- The Learning can be done without examples of optimal behaviour.
- Feedback is delayed, not instantaneous.
- Time really matters (sequential, non i.i.d data).
- Agent's actions affect the subsequent data it receives.



Examples of Reinforcement Learning



- Fly manoeuvres in a helicopter
- Manage an investment portfolio
- Control a power station
- Make a humanoid robot walk
- Play many different video games better than humans

Helicopter Manoeuvres



- <http://heli.stanford.edu/>

Paper: [Autonomous Helicopter Aerobatics through Apprenticeship Learning](#), Pieter Abbeel, Adam Coates, and Andrew Y. Ng, 2010

Week 8: Reinforcement Learning

About RL


Traffic Light Control



Week 8: Reinforcement Learning

About RL

Make a robot walk



<https://www.dropbox.com/s/fdn1loibsh2p0sa/parkour.mp4?e=1&dl=0>

Rewards



- A **reward** R_t is a scalar feedback signal
- Indicates how well agent is doing at step t
- The agent's job is to maximise cumulative reward

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} + \dots$$

- We call this the **return**

Reinforcement learning is based on the **reward hypothesis**.

Reward Hypothesis

All goals can be described by the maximisation of expected cumulative reward

Examples of Rewards



- Fly manoeuvres in a helicopter
 - + reward for following desired trajectory
 - - reward for crashing
- Manage an investment portfolio
 - + reward for each \$ in bank
- Control a power station
 - + reward for producing power
 - - reward for exceeding safety thresholds
- Make a humanoid robot walk
 - + reward for forward motion
 - - reward for falling over
- Play many different video games better than humans
 - +/- reward for increasing/decreasing score

Sequential Decision Making



- Goal: select actions to maximise total future reward
- Actions may have long term consequences
- Reward may be delayed
- It may be better to sacrifice immediate reward to gain more long-term reward.
- Examples:
 - A financial investment (may take months to mature)
 - Refuelling a helicopter (might prevent a crash in several hours)
 - Blocking opponent moves (might help winning chances many moves from now)


Week 8: Reinforcement Learning

The RL Problem: Environments

Agent and Environment



- At each step t the agent:
 - Receives Observation O_t
 - Executes Action A_t
 - Receives Reward R_t
- The environment
 - Receives action A_t
 - Emits observation O_{t+1}
 - Emits reward R_{t+1}
- t increments at env. step



Ref: RL course by David Silver

History and state



- The **history** is the full sequence of observations, actions, rewards

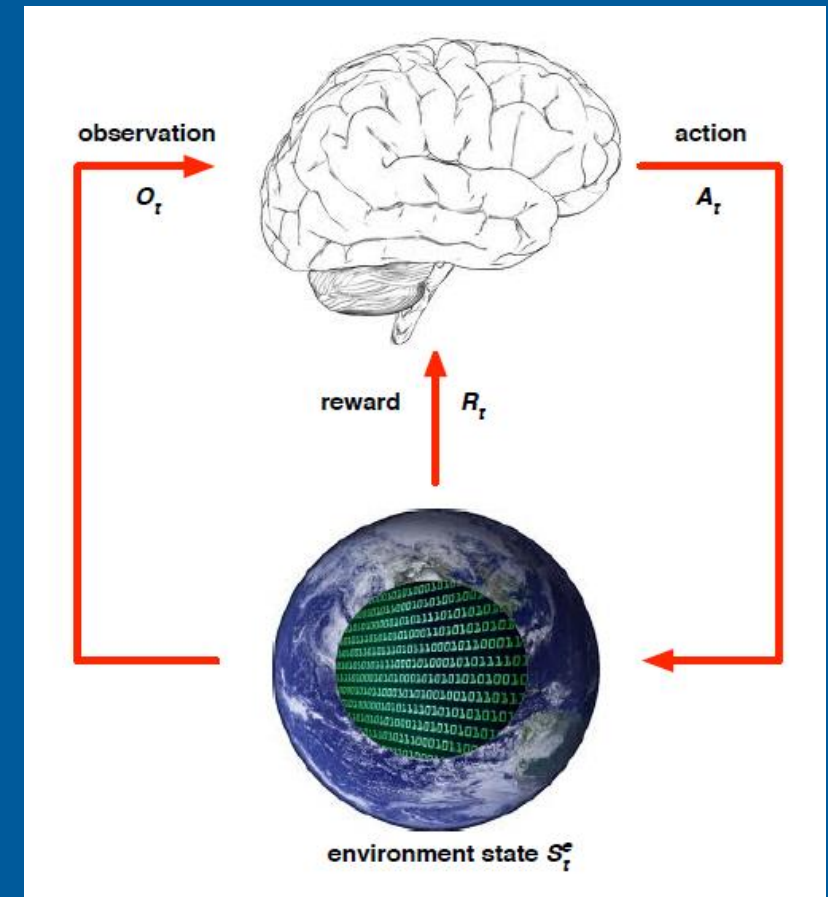
$$H_t = O_0, A_0, R_1, O_1, \dots, O_{t-1}, A_{t-1}, R_t, O_t$$

- i.e. all observable variables up to time t
- **State** is the information used to determine what happens next
- Formally, state is a function of the history:

$$S_t = f(H_t)$$

Environment state

- The **environment state** is the environment's internal state
- It is usually invisible to the agent
- Even if it is visible, it may contain lots of irrelevant information

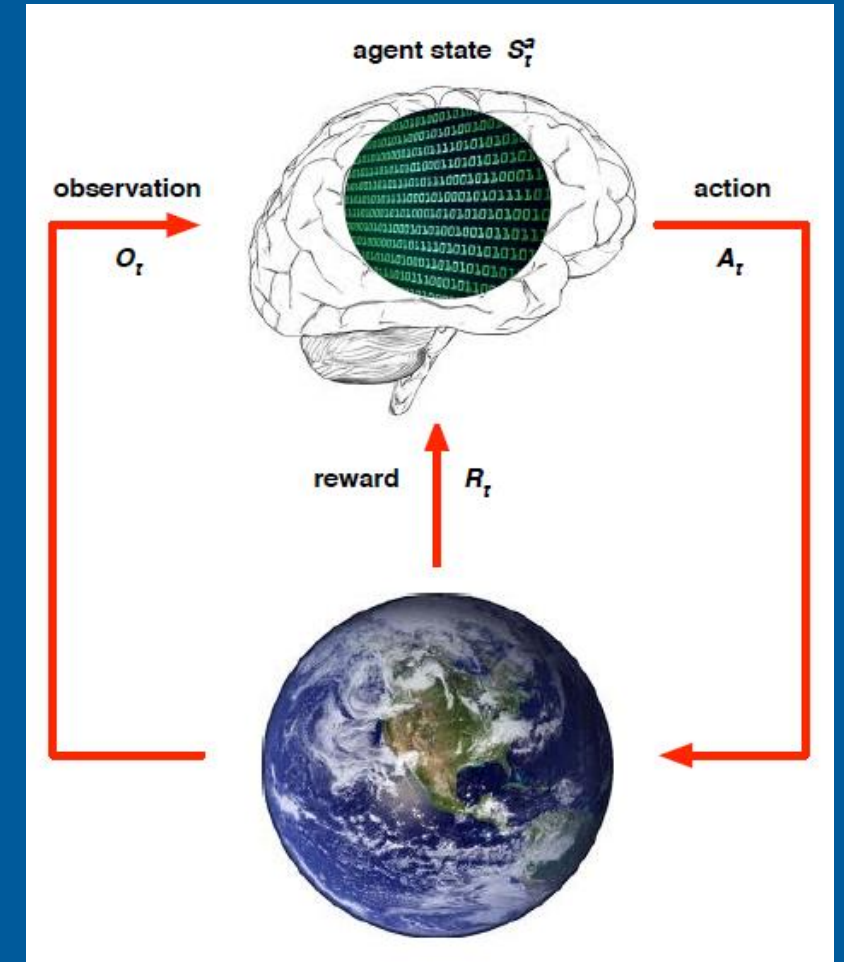


Ref: RL course by David Silver

Agent state

- The agent state S_t is the agent's internal representation
- i.e. whatever information the agent uses to pick the next action
- i.e. it is the information used by reinforcement learning algorithms
- **Agent State** is the information used to determine what happens next
- It can be any function of history:

$$S_t = f(H_t)$$



Ref: RL course by David Silver

Information state



- An **information state** (a.k.a. **Markov state**) contains all useful information from the history.

Definition

A state S_t is Markov if and only if


$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, \dots, S_t]$$

- The future is independent of the past given the present
- Once the state is known, the history may be thrown away
- i.e. The state is a sufficient statistic of the future
- The history H_t is Markov.

Week 8: Reinforcement Learning

The RL Problem: state

Rat Example




Ref: RL course by David Silver

- What if agent state = last 3 items in sequence?
- What if agent state = counts for lights, bells and levers?
- What if agent state = complete sequence?

Fully observable environment

- Agent directly observes environment state
- Agent state= environment state= information state
$$S_t = O_t = \text{environment state}$$
- Formally, this is a Markov decision process (MDP)



Ref: RL course by David Silver

Partially Observable Environment



- **Partial observability:** agent **indirectly** observes environment:
 - A robot with camera vision isn't told its absolute location
 - A poker playing agent only observes public cards
- Now agent state \neq environment state
- using the observation as state would not be Markovian
- Formally this is a **partially observable Markov decision process** (POMDP)

Major Components of an RL Agent



- An RL agent may include one or more of these components:
 - **Policy**: agent's behaviour function
 - **Value function**: how good is each state and/or action
 - **Model**: agent's representation of the environment

Week 8: Reinforcement Learning

Inside An RL Agent

Policy



- A **policy** defines the agent's behaviour
- It is a map from state to action
- Deterministic policy: $a = \pi(s)$
- Stochastic policy: $\pi(a|s) = \mathbb{P}(A_t = a|S_t = s)$

Value Function



- **Value function** is a prediction of future reward
- Can be used to evaluate the goodness/badness of states
- Can be used to select between actions
- The actual value function is the expected return

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s] = \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$

- Discount factor $\gamma \in [0,1]$: Trades off importance of immediate vs long-term rewards

Week 8: Reinforcement Learning

Inside An RL Agent

Model



- A **model** predicts what the environment will do next
- Predicts the next state

$$\mathcal{P}_{s\acute{s}}^a = \mathbb{P}[S_{t+1} = \acute{s} | S_t = s, A_t = a]$$

- Predicts the next reward

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

- A model does not immediately give us a good policy - we would still need to plan


Week 8: Reinforcement Learning Inside An RL Agent

Maze Example

- Rewards: -1 per time-step
- Actions: N, E, S, W
- States: Agent's location



Start




Goal

Ref: RL course by David Silver

Week 8: Reinforcement Learning
Inside An RL Agent

Maze Example: Policy

- Arrows represent policy $\pi(s)$ for each state s




Ref: RL course by David Silver

Maze Example: Value function

- Numbers represent value $v_{\pi}(s)$ of each state s



Start




Goal

Ref: RL course by David Silver

Maze Example: Model

- Agent may have an internal model of the environment
- Dynamics: how actions change the state
- Rewards: how much reward from each state
- Grid layout represents transition model $\mathcal{P}_{s\acute{s}}^a$
- Numbers represent immediate reward \mathcal{R}_s^a from each state s (same for all a and \acute{s} in this case)

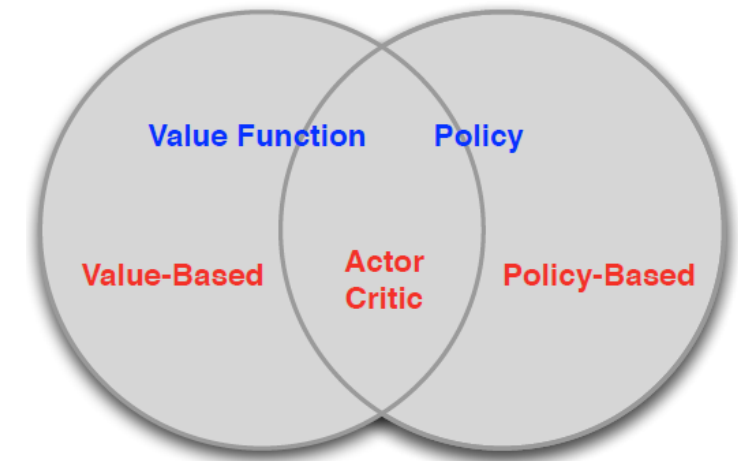


Ref: RL course by David Silver

Categorizing RL agents(1): Value based and Policy based



- Value- based : will determine a value function that quantifies the reward and using this value function we determine the optimal policy
 - No policy(implicit)
 - Value function
 - Q- learning
 - Deep Q network
 - SARSA
- Policy based: will determine an optimal policy directly which means the policy that maximizes reward
 - Policy
 - No value function
 - REINFORCE
 - PPO (Proximal Policy Optimization)
 - TRPO (Trust Region Policy Optimization)
- Actor Critic
 - Policy network (Actor)
 - Value function (Critic)



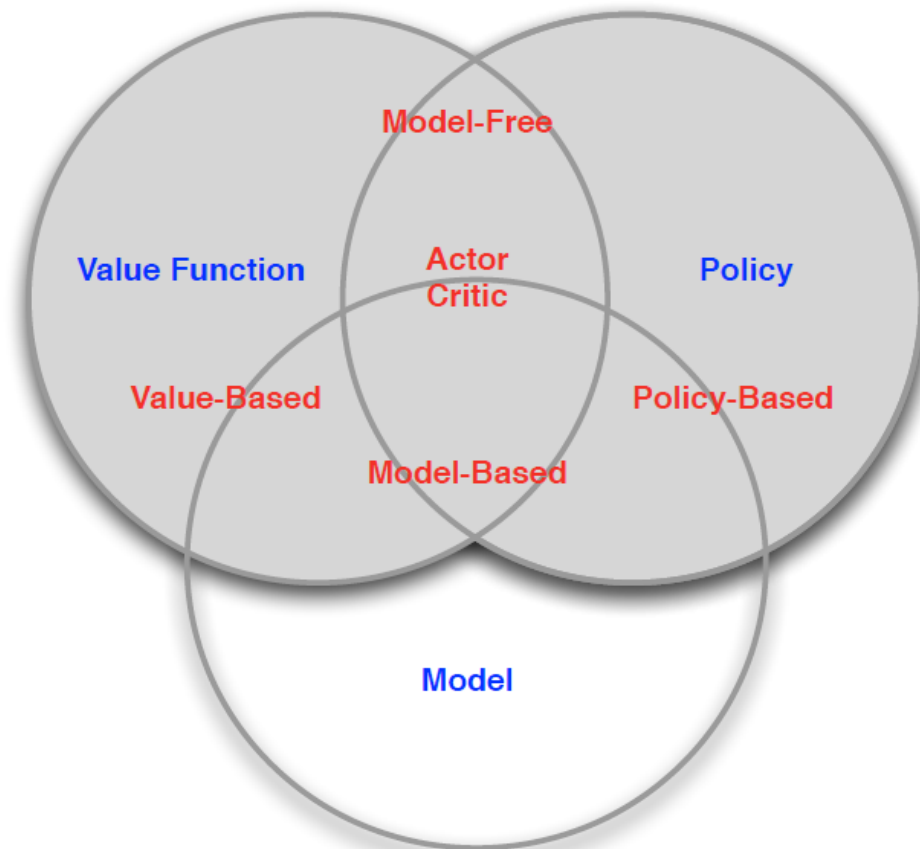
Categorizing RL agents(2): Model free and Model Based



- Model free
 - Policy and/or Value Function
 - No Model
- Model based
 - Policy and/or Value Function
 - Model

Week 8: Reinforcement Learning Inside An RL Agent

RL Agent taxonomy



Ref: RL course by David Silver


Summary: Key concepts in RL



- **Environment:** Physical world in which the agent operates
- **State:** Current situation of the agent
- **Reward:** Positive or negative feedback from the environment
- **Policy:** The rules that change agent's state to actions
- **Value:** Future reward that an agent would receive

Week 8: Reinforcement Learning
Problems within RL

Value Based Methods




Week 8: Reinforcement Learning
Problems within RL

Value Functions




State- value functions $V(s)$



- How good is it to be in the state s

State-action value functions $Q(s, a)$



- How good is it to be in the state s and take an action a in this state

Bellman Equation



- fundamental concept in dynamic programming and reinforcement learning, named after Richard Bellman, who introduced it in the 1950s
 - It provides a recursive decomposition for solving optimization problems, particularly those involving decision-making over time.
 - the Bellman equation is used to describe the relationship between the value of a state and the values of subsequent states.
 - Bellman equation provides a way to compute the value of each state (or state-action pair) recursively by considering the expected rewards and the values of subsequent states.
-
- Bellman Expectation Equation
 - Bellman optimality Equation

Bellman Expectation Equations



- For value function

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma V_{\pi}(S_{t+1}) \mid S_t = s]$$

- $V_{\pi}(s)$ is the value of state s under policy π
- \mathbb{E}_{π} is the expected value given that the agent follows policy π
- R_{t+1} is the reward received after transitioning from state S to state S_{t+1}

- For Q-values

$$Q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma Q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

Bellman Optimality Equations



■ For value function

- For the optimal value function $V^*(s)$, which represents the maximum expected return achievable from state s , the Bellman optimality equation is:

$$V^*(s) = \max_a \mathbb{E}[R_{t+1} + \gamma V^*(S_{t+1}) \mid S_t = s, A_t = a]$$

- $V^*(s)$ is the optimal value of state s
- \max_a is the maximization over all possible actions a .
- The expectation \mathbb{E} is taken over the possible next states S_{t+1} and rewards R_{t+1} , given action $A_t = a$ in state s

■ For Q-values

$$Q^*(s, a) = \mathbb{E}[R_{t+1} + \gamma \max_{\hat{a}}(S_{t+1}, \hat{a}) \mid S_t = s, A_t = a]$$

Exploration and Exploitation



- Reinforcement learning is like trial-and-error learning
- The agent should discover a good policy
- From its experiences of the environment
- Without losing too much reward along the way
- **Exploration** finds more information about the environment
 - trying out new actions that may not be the best according to the agent's current knowledge, but could potentially lead to discovering better long-term strategies.
- **Exploitation** exploits known information to maximise reward
 - The agent uses its current knowledge to choose actions that it believes will give the highest reward based on past experiences.
- It is usually important to explore as well as exploit

Week 8: Reinforcement Learning

Problems within RL

Examples



- Restaurant Selection

Exploitation Go to your favourite restaurant

Exploration Go to new restaurant

- Online Banner Advertisements

Exploitation Show the most successful advert

Exploration Show a different advert

- Oil Drilling

Exploitation Drill at the best known location

Exploration Drill at a new location

- Game Playing

Exploitation Play the move you believe is best

Exploration Play an experimental move

The Exploration-Exploitation Trade-off



- RL agent needs to make decisions about whether to use known strategies to get immediate rewards (exploitation) or try new strategies that might lead to better rewards in the future (exploration).
- **Too much exploitation:** The agent might not find better strategies that could improve its performance in the long run, resulting in less effective outcomes over time.
- **Too much exploration:** The agent may spend too much time trying new actions, resulting in lower immediate rewards and slow learning.
- Balancing these two approaches is known as the exploration-exploitation trade-off.
- **Epsilon-Greedy Strategy:**
 - select a random action with probability ϵ (exploration)
 - Select the best-known action with probability $1-\epsilon$ (exploitation). $a = \operatorname{argmax}_{a \in A} Q(a)$
 - Lower ϵ over time: Often, ϵ starts high to encourage exploration and gradually decreases to shift towards more exploitation.

Learning and Planning



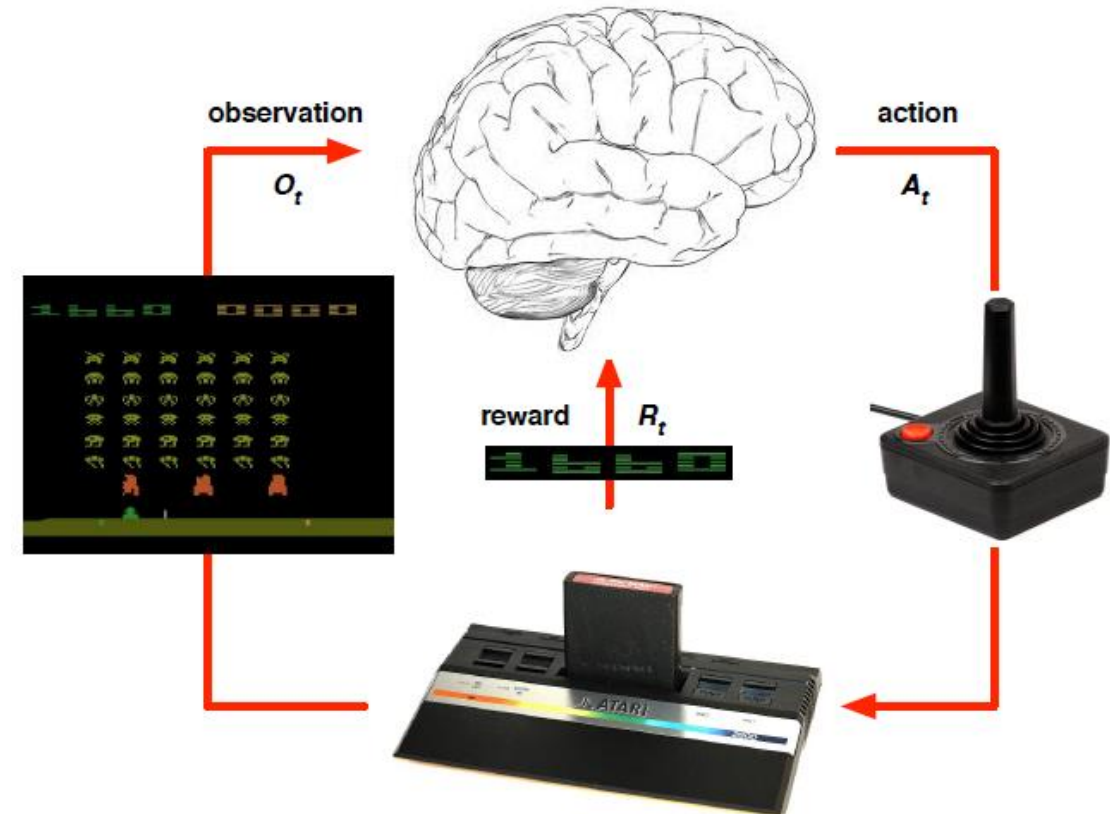
Two fundamental problems in sequential decision making

- Reinforcement Learning
 - The environment is initially unknown
 - The agent interacts with the environment
 - The agent improves its policy
- Planning
 - A model of the environment is known
 - The agent performs computations with its model (without any external interaction)
 - The agent improves its policy
 - a.k.a. reasoning, thought, search, planning

Atari Example: Reinforcement Learning



- Rules of the game are unknown
- Learn directly from interactive game-play
- Pick actions on joystick, see pixels and scores

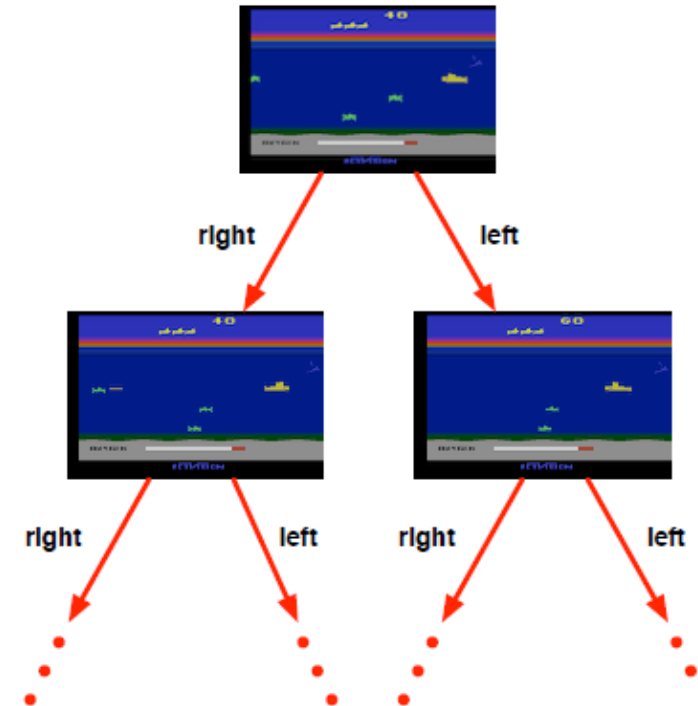


Ref: RL course by David Silver

Atari Example: Planning



- Rules of the game are known
- Can query emulator
 - perfect model inside agent's brain
- If I take action a from state s :
 - what would the next state be?
 - what would the score be?
- Plan ahead to find optimal policy
 - e.g. tree search



Ref: RL course by David Silver