



Technische Hochschule  
Ingolstadt

# Principles of Autonomy and Decision Making

(AI\_PrincAutonomy\_2808)

Week 9: Deep Dive into Reinforcement Learning  
and Q-Learning

Team:

- Prof. Dr. rer. nat. Lenz Belzner
- Chidvilas Karpenahalli Ramakrishna, M.Eng.
- Adithya Mohan, M.Eng.
- Zahra Zeinaly, Ph.D.



# Announcements

- **Assignment 1: Build your own environment**
  - **Deadline:** Monday, 10 June 2024, 23:59
  - **Rules:**
    - Please do not use Dynamic Environments!
    - **If you are not using images:** You should submit one `.py` file with the `"env_<student user id>.py"` naming convention. For example, `"env_chk3541.py"`
    - **If you are using images:** Please submit the assignment as one `.zip` file with the naming convention `"env_<student user id>.zip"`
- **Assignment 2: MDPs and Value Iteration**
  - **Deadline:** Monday, 10 June 2024, 23:59
- **Assignment 3: Build and train a Q-Learning agent**
  - Coming soon!
  - **Open date:** Thursday, 06.06.2024, 18:00
  - **Deadline:** Sunday, 30.06.2024, 23:59
- **Final Presentation:**
  - **Tentative dates:** 01.07.2024 to 10.07.2024
  - Slots will be available soon on Moodle. Each student can only book one slot.
  - Slots fixed and are on First-Come-First-Serve basis
  - *Slots are binding! Cancellations are not allowed!*

# Week 9: Deep Dive into Reinforcement Learning and Q-Learning

## Recap



Complexity



### Uninformed Search

- Depth-First Search (DFS)
  - No cost consideration
- Breadth-First Search (BFS)
  - No cost consideration
- Uniform Cost Search (UCS)
  - Cost consideration

### Informed Search

- Greedy Search
  - Heuristics
- A\* algorithm
  - Cost + Heuristics

### MDP

- Uncertainty consideration
  - $T(s, a, s')$
- Reward consideration
  - $R(s, a, s')$

Value Iteration

Policy Iteration

### POMDP

- Uncertainty consideration
  - $T(s, a, s')$
  - $O(o, s', a)$
- Reward consideration
  - $R(s, a, s')$

MC Methods and  
Particle Filtering

### Question:

- Do we have access to  $T(s, a, s')$ ,  $R(s, a, s')$  and  $O(o, s', a)$ ?

## Why Reinforcement Learning?



- MDP Bellman equation:

$$V^*(s) = \max_a \sum_{s'} P(s'|a,s) [ R(s,a,s') + \gamma V^*(s') ]$$

- POMDP Bellman equation:

$$V(b) = \max_{a \in A} \left[ \sum_{s \in S} b(s) \left\{ \sum_{s' \in S} P(s'|s,a) \left( R(s,a,s') + \gamma \sum_{o \in O} P(o|s',a) V(b') \right) \right\} \right]$$

- In reality:

- We may not have access to the set of states in the environment i.e.  $S = \{s_1, s_2, \dots, s_n\}$
- We may not have access to  $P(s'|a,s)$
- We may not have access to  $R(s,a,s')$
- This would result in...

$$V^*(s) = \max_a \sum_{s'} \text{🐒} [ \text{🐒} + \gamma V^*(s') ] \text{ i.e. Unknown MDP}$$

- How would you solve it?

➡ Reinforcement Learning (RL) 😊

- Solution:

- We have to explore the environment to find states
- We have to interact with the environment to learn  $P(s'|a,s)$  and  $R(s,a,s')$

## Model-based and Model-free approach



- Example:
  - Estimate the average test score of PADM students
- Model-based approach:
  - I have data from previous tests showing how different study habits that influence the test scores like hours spent studying, types of study materials used etc
  - I create a model to predict the test scores based on these study habits
  - I use this model to predict the expected test score for each student and calculate the average score
- Model-Free Approach:
  - I don't have data on study habits or a predictive model
  - I collect the scores of a sample of students who take a mock test
  - I use the frequency of these scores to calculate the average score
- In the current lecture, we consider model-free approaches

# Unknown MDP



### Known MDP:

- $T(s, a, s') = 70\%$  in intended direction
- $T(s, a, s') = 10\%$  in each of the unintended direction





Fig. 1 A grid-world with known MDP

### Unknown MDP:

- $T(s, a, s') = ?$
- We will gather experiences and predict the transition probability i.e.  $\hat{T}(s, a, s')$
- $T(s, a, s')$  is built into the sampling process



Outcome  
→  $\hat{T}(s, a, s') = 60\%$  in intended direction

In RL, we learn the MDP

Fig. 2 A grid-world with unknown MDP

## Background: Moving average



- We have to update our knowledge as we experience new things

- Cumulative Moving Average (CMA):

- CMA considers all the data points up to the current point

$$CMA_t = \frac{1}{t} \sum_{i=1}^t x_i$$

- Simple Moving Average (SMA):

- SMA is calculated by taking the arithmetic mean of a fixed number of recent values in the dataset

$$SMA_t = \frac{1}{n} \sum_{i=t-n+1}^t x_i$$

- Exponential moving average (EMA):

- EMA gives more weight to recent data points, making it more responsive to new information compared to the SMA

$$EMA_t = (1 - \alpha).EMA_{t-1} + \alpha.x_t$$

## Q-Learning



- Value-Iteration:

$$V_k(s) = \max_a \sum_{s'} P(s'|a, s) [ R(s, a, s') + \gamma V_{k-1}(s') ]$$

- Q-value iteration:

- We are more interested in the Q-values
- Initialize all Q-values to 0 i.e.  $Q_k(s, a) = 0$
- Updating is similar to value-iteration

$$Q_k(s, a) = \sum_{s'} P(s'|a, s) [ R(s, a, s') + \gamma \max_{a'} Q_{k-1}(s', a') ]$$

- However, we do not know  $R(s, a, s')$  and  $P(s'|a, s)$
- Solution:
  - Q-learning



## Week 9: Deep Dive into Reinforcement Learning and Q-Learning

# Q-Learning



- Q-learning:
  - Use moving averages i.e. compute averages as we go
  - Gather a sample by interacting with the environment i.e.  $(s, a, s', r)$ 
    - $T(s, a, s')$  governs this transition
  - $sample = [R(s, a, s') + \gamma \max_{a'} Q_{k-1}(s', a')]$  by considering the old Q-values i.e.  $Q_{k-1}(s', a')$
  - $Q_k(s, a) = (1 - \alpha)Q_{k-1}(s, a) + \alpha \cdot (sample)$
- After substitution and rearrangement:

$$Q_k(s, a) = Q_{k-1}(s, a) + \alpha \cdot [R(s, a, s') + \gamma \max_{a'} Q_{k-1}(s', a') - Q_{k-1}(s, a)] \longrightarrow \text{Q-learning update rule}$$

Diagram illustrating the Q-learning update rule with annotations:

- New Q-value (points to  $Q_k(s, a)$ )
- Old Q-value (points to  $Q_{k-1}(s, a)$ )
- Learning rate (points to  $\alpha$ )
- Immediate reward for transitioning into  $s'$  (points to  $R(s, a, s')$ )
- Discount factor (points to  $\gamma$ )
- Maximum expected future reward (points to  $\max_{a'} Q_{k-1}(s', a')$ )
- Old Q-value (points to  $Q_{k-1}(s, a)$ )

## Features of Q-Learning



- Q-learning is a “model-free” algorithm
  - Does not require the model of the environment
  - It just needs the ability to sample transitions and rewards
- Q-learning is an “off-policy” algorithm
  - Q-learning is called off-policy because it decouples the learning of the optimal policy from the behaviour policy used to explore the environment
  - This allows Q-learning to effectively learn the optimal action-value function, even when the agent's actions are driven by an exploratory or suboptimal policy
- Q-learning convergence to optimal policy
  - Conditions:
    1. The agent must “sufficiently” explore
    2. The learning rate has to decay over time
    3. The environment must remain static/ stationary i.e.  $T(s, a, s')$  and  $R(s, a, s')$  should remain same

## Exploration vs. Exploitation



- Exploration:
  - This involves selecting an action at random, which allows the agent to discover new states and potentially learn more about the environment, even if it means potentially receiving a lower immediate reward
- Exploitation:
  - This involves selecting the action that has the highest known Q-value for a given state, effectively leveraging the knowledge already gathered to maximize immediate rewards
- Q-learning uses an  $\epsilon$ -greedy policy for exploration
  - $\epsilon \in [0, 1]$
  - $\epsilon = 0$  means the agent always exploits and  $\epsilon = 1$  means the agent only explores
  - So, the agent chooses exploration with a probability of  $\epsilon$  and chooses the action with the highest Q-value for the given state with a probability of  $(1 - \epsilon)$

## Week 9: Deep Dive into Reinforcement Learning and Q-Learning

# Q-table



- A Q-table is used to store and update Q-values as shown in Fig. 1
- New rows are added as the agent discovers new states
- The Q-values of each state-action pair is updated as per the Q-learning update rule
- Once training is done, when the agent visits the state  $s_n$ , it chooses action based on the maximum Q-value
- Can you see the drawbacks of Q-learning?

Actions States	Action 1 „Up“	Action 2 „Down“	Action 3 „Left“	Action 4 „Right“
State 1 $s_1$	$Q(s_1, U)$	$Q(s_1, D)$	$Q(s_1, L)$	$Q(s_1, R)$
State 2 $s_2$	$Q(s_2, U)$	$Q(s_2, D)$	$Q(s_2, L)$	$Q(s_2, R)$
.	.	.	.	.
.	.	.	.	.
.	.	.	.	.
State n $s_n$	$Q(s_n, U)$	$Q(s_n, D)$	$Q(s_n, L)$	$Q(s_n, R)$

Fig. 1 Q-table with max Q-values highlighted in green

## Drawbacks of Q-learning



- Scalability issues:
  - Not suitable for high dimensional state spaces – more rows
  - Not suitable for high dimensional action spaces – more columns
  - Not suitable for continuous action and state spaces
- Non-static environment: Cannot handle complex state spaces i.e. no approximation capabilities
  - Do not use a Q-learning to solve a Dynamic environment
  - If you wish to use Q-learning for a Dynamic environment, you need „Approximate Q-learning“
  - „Approximate Q-learning“ is out of scope of the current lecture